

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
```

In [2]:

```
1 df=pd.read_csv(r"C:\Users\HP\OneDrive\Documents\used_cars_data.csv")
2 df
```

Out[2]:

|      | S.No. | Name  | Location   | Year | Kilometers_Driven | Fuel_Type | Transmission | Ow  |
|------|-------|---|------------|------|-------------------|-----------|--------------|-----|
| 0    | 0     | Maruti Wagon R LXI CNG                            | Mumbai     | 2010 | 72000             | CNG       | Manual       |     |
| 1    | 1     | Hyundai Creta 1.6 CRDi SX Option                  | Pune       | 2015 | 41000             | Diesel    | Manual       |     |
| 2    | 2     | Honda Jazz V                                      | Chennai    | 2011 | 46000             | Petrol    | Manual       |     |
| 3    | 3     | Maruti Ertiga VDI                                 | Chennai    | 2012 | 87000             | Diesel    | Manual       |     |
| 4    | 4     | Audi A4 New 2.0 TDI Multitronic                   | Coimbatore | 2013 | 40670             | Diesel    | Automatic    |     |
| ...  | ...   | ...   | ...        | ...  | ...               | ...       | ...          | ... |
| 7248 | 7248  | Volkswagen Vento Diesel Trendline                 | Hyderabad  | 2011 | 89411             | Diesel    | Manual       |     |
| 7249 | 7249  | Volkswagen Polo GT TSI                            | Mumbai     | 2015 | 59000             | Petrol    | Automatic    |     |
| 7250 | 7250  | Nissan Micra Diesel XV                            | Kolkata    | 2012 | 28000             | Diesel    | Manual       |     |
| 7251 | 7251  | Volkswagen Polo GT TSI                            | Pune       | 2013 | 52262             | Petrol    | Automatic    |     |
| 7252 | 7252  | Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan... | Kochi      | 2014 | 72443             | Diesel    | Automatic    |     |

7253 rows × 14 columns



In [3]:


```
1 df = df[['Kilometers_Driven', 'Price']]
2 df.columns=['Km', 'Cost']
```

In [4]:

```
1 sns.lmplot(x='Km',y='Cost',data=df,order=2,ci=None)
```

Out[4]:

&lt;seaborn.axisgrid.FacetGrid at 0x201d0f19930&gt;



In [5]:

```
1 df.describe()
```

Out[5]:

|       | Km           | Cost        |
|-------|--------------|-------------|
| count | 7.253000e+03 | 6019.000000 |
| mean  | 5.869906e+04 | 9.479468    |
| std   | 8.442772e+04 | 11.187917   |
| min   | 1.710000e+02 | 0.440000    |
| 25%   | 3.400000e+04 | 3.500000    |
| 50%   | 5.341600e+04 | 5.640000    |
| 75%   | 7.300000e+04 | 9.950000    |
| max   | 6.500000e+06 | 160.000000  |

In [6]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Km      7253 non-null    int64  
 1   Cost    6019 non-null    float64
dtypes: float64(1), int64(1)
memory usage: 113.5 KB
```

In [7]:

```
1 df.fillna(method='ffill',inplace=True)
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_8628\4116506308.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.fillna(method='ffill',inplace=True)
```

In [8]:

```
1 df.dropna(inplace=True)
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_8628\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.dropna(inplace=True)
```

In [9]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7253 entries, 0 to 7252
```

```
Data columns (total 2 columns):
```

| # | Column | Non-Null Count | Dtype   |
|---|--------|----------------|---------|
| 0 | Km     | 7253 non-null  | int64   |
| 1 | Cost   | 7253 non-null  | float64 |

```
dtypes: float64(1), int64(1)
```

```
memory usage: 113.5 KB
```

In [10]:

```
1 df.isnull().sum()
```

Out[10]:

```
Km      0
```

```
Cost     0
```

```
dtype: int64
```

In [11]:

```
1 df.head(10)
```

Out[11]:


|   | Km    | Cost  |
|---|-------|-------|
| 0 | 72000 | 1.75  |
| 1 | 41000 | 12.50 |
| 2 | 46000 | 4.50  |
| 3 | 87000 | 6.00  |
| 4 | 40670 | 17.74 |
| 5 | 75000 | 2.35  |
| 6 | 86999 | 3.50  |
| 7 | 36000 | 17.50 |
| 8 | 64430 | 5.20  |
| 9 | 65932 | 1.95  |

In [12]:

```
1 sns.lmplot(x='Km',y='Cost',data=df,order=2,ci=None)
```

Out[12]:

&lt;seaborn.axisgrid.FacetGrid at 0x201b6d18370&gt;



In [13]:

```
1 x=np.array(df['Km']).reshape(-1,1)
2 y=np.array(df['Cost']).reshape(-1,1)
```

In [14]:

```
1 x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.25)
2 regr = LinearRegression()
3 regr.fit(x_train,y_train)
4 print(regr.score(x_test,y_test))
```


-0.0009939448569544762

In [15]:

```
1 y_pred = regr.predict(x_test)
2 plt.scatter(x_test,y_test,color='b')
3 plt.plot(x_test,y_pred,color='c')
```

Out[15]:

[<matplotlib.lines.Line2D at 0x201b6e012d0>]



In [16]:

```
1 df200 = df[:, :200]
2 df200
```

Out[16]:


|     | Km     | Cost  |
|-----|--------|-------|
| 0   | 72000  | 1.75  |
| 1   | 41000  | 12.50 |
| 2   | 46000  | 4.50  |
| 3   | 87000  | 6.00  |
| 4   | 40670  | 17.74 |
| ... | ...    | ...   |
| 195 | 52000  | 3.50  |
| 196 | 43571  | 3.55  |
| 197 | 50000  | 3.25  |
| 198 | 113000 | 4.50  |
| 199 | 90000  | 5.35  |

200 rows × 2 columns

In [17]:

```
1 df200.fillna(method='ffill',inplace=True)
2 X=np.array(df['Km']).reshape(-1,1)
3 y=np.array(df['Cost']).reshape(-1,1)
4 df200.dropna(inplace=True)
5 X_train,x_test,y_train,y_test = train_test_split(X,y,test_size=0.25)
6 regr=LinearRegression()
7 regr.fit(X_train,y_train)
8 print("Regressin: ",regr.score(x_test,y_test))
9 y_pred=regr.predict(x_test)
10 plt.scatter(x_test,y_test,color='b')
11 plt.plot(x_test,y_pred,color='k')
12 plt.show()
```

Regressin: -0.0017436033723474686



In [18]:

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model = LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred = model.predict(x_test)
6 r2=r2_score(y_test,y_pred)
7 print('R2 score: ',r2)
```

R2 score: -0.0017436033723474686



In [19]:

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn import metrics
3 model = LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred = model.predict(x_test)
6 r2=metrics.mean_squared_error(y_test,y_pred)
7 print('MSE: ',r2)
8
```

MSE: 104.74884224431655

## Logistic Regression

In [20]:

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.preprocessing import StandardScaler
```

In [21]:

```
1 a=pd.read_csv(r"C:\Users\HP\OneDrive\Documents\used_cars_data-2.csv")
2 a
```

Out[21]:

|      | S.No. | Name  | Location   | Year | Kilometers_Driven | Fuel_Type | Transmission | Mileage |
|------|-------|---|------------|------|-------------------|-----------|--------------|---------|
| 0    | 0     | Maruti Wagon R LXI CNG                            | Mumbai     | 2010 | 72000             | CNG       | Manual       | k       |
| 1    | 1     | Hyundai Creta 1.6 CRDi SX Option                  | Pune       | 2015 | 41000             | Diesel    | Manual       | 1       |
| 2    | 2     | Honda Jazz V                                      | Chennai    | 2011 | 46000             | Petrol    | Manual       |         |
| 3    | 3     | Maruti Ertiga VDI                                 | Chennai    | 2012 | 87000             | Diesel    | Manual       | 2       |
| 4    | 4     | Audi A4 New 2.0 TDI Multitronic                   | Coimbatore | 2013 | 40670             | Diesel    | Automatic    |         |
| ...  | ...   | ...   | ...        | ...  | ...               | ...       | ...          |         |
| 7248 | 7248  | Volkswagen Vento Diesel Trendline                 | Hyderabad  | 2011 | 89411             | Diesel    | Manual       | 2       |
| 7249 | 7249  | Volkswagen Polo GT TSI                            | Mumbai     | 2015 | 59000             | Petrol    | Automatic    | 1       |
| 7250 | 7250  | Nissan Micra Diesel XV                            | Kolkata    | 2012 | 28000             | Diesel    | Manual       | 2       |
| 7251 | 7251  | Volkswagen Polo GT TSI                            | Pune       | 2013 | 52262             | Petrol    | Automatic    |         |
| 7252 | 7252  | Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan... | Kochi      | 2014 | 72443             | Diesel    | Automatic    |         |

7253 rows × 14 columns



In [22]:

```
1 a.describe()
```

Out[22]:

|       | S.No.       | Year        | Kilometers_Driven | Seats       | Price       |
|-------|-------------|-------------|-------------------|-------------|-------------|
| count | 7253.000000 | 7253.000000 | 7.253000e+03      | 7200.000000 | 6019.000000 |
| mean  | 3626.000000 | 2013.365366 | 5.869906e+04      | 5.279722    | 9.479468    |
| std   | 2093.905084 | 3.254421    | 8.442772e+04      | 0.811660    | 11.187917   |
| min   | 0.000000    | 1996.000000 | 1.710000e+02      | 0.000000    | 0.440000    |
| 25%   | 1813.000000 | 2011.000000 | 3.400000e+04      | 5.000000    | 3.500000    |
| 50%   | 3626.000000 | 2014.000000 | 5.341600e+04      | 5.000000    | 5.640000    |
| 75%   | 5439.000000 | 2016.000000 | 7.300000e+04      | 5.000000    | 9.950000    |
| max   | 7252.000000 | 2019.000000 | 6.500000e+06      | 10.000000   | 160.000000  |

In [23]:

```
1 a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   S.No.                  7253 non-null  int64
1   Name                   7253 non-null  object
2   Location               7253 non-null  object
3   Year                   7253 non-null  int64
4   Kilometers_Driven      7253 non-null  int64
5   Fuel_Type              7253 non-null  object
6   Transmission           7253 non-null  object
7   Mileage                7251 non-null  object
8   Engine                 7207 non-null  object
9   Power                  7207 non-null  object
10  Seats                  7200 non-null  float64
11  New_Price              1006 non-null  object
12  Price                  6019 non-null  float64
13  Owner_Type             7253 non-null  object
dtypes: float64(2), int64(3), object(9)
memory usage: 793.4+ KB
```

In [24]:

```
1 a.isnull().sum()
```

Out[24]:

```
S.No.          0
Name           0
Location       0
Year           0
Kilometers_Driven  0
Fuel_Type      0
Transmission   0
Mileage        2
Engine         46
Power          46
Seats          53
New_Price      6247
Price          1234
Owner_Type     0
dtype: int64
```

In [25]:

```
1 a.fillna(method='ffill',inplace=True)
```

In [26]:

```
1 a.dropna(inplace=True)
```

In [27]:

```
1 a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 7251 entries, 2 to 7252
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   S.No.                 7251 non-null  int64
1   Name                  7251 non-null  object
2   Location              7251 non-null  object
3   Year                  7251 non-null  int64
4   Kilometers_Driven     7251 non-null  int64
5   Fuel_Type             7251 non-null  object
6   Transmission          7251 non-null  object
7   Mileage               7251 non-null  object
8   Engine                7251 non-null  object
9   Power                 7251 non-null  object
10  Seats                 7251 non-null  float64
11  New_Price             7251 non-null  object
12  Price                 7251 non-null  float64
13  Owner_Type            7251 non-null  object
dtypes: float64(2), int64(3), object(9)
memory usage: 849.7+ KB
```

In [28]:

```
1 a.isnull().sum()
```

Out[28]:

```
S.No.      0
Name        0
Location    0
Year        0
Kilometers_Driven  0
Fuel_Type   0
Transmission  0
Mileage      0
Engine      0
Power       0
Seats       0
New_Price   0
Price       0
Owner_Type  0
dtype: int64
```

In [29]:

```
1 print("This DataFrame has %d rows and %d columns"%(a.shape))
```

This DataFrame has 7251 rows and 14 columns

In [30]:

```
1 a.head()
```

Out[30]:

|   | S.No. | Name                            | Location   | Year | Kilometers_Driven | Fuel_Type | Transmission | Mileage    |
|---|-------|---------------------------------|------------|------|-------------------|-----------|--------------|------------|
| 2 | 2     | Honda Jazz V                    | Chennai    | 2011 | 46000             | Petrol    | Manual       | 18.2 kmpl  |
| 3 | 3     | Maruti Ertiga VDI               | Chennai    | 2012 | 87000             | Diesel    | Manual       | 20.77 kmpl |
| 4 | 4     | Audi A4 New 2.0 TDI Multitronic | Coimbatore | 2013 | 40670             | Diesel    | Automatic    | 15.2 kmpl  |
| 5 | 5     | Hyundai EON LPG Era Plus Option | Hyderabad  | 2012 | 75000             | LPG       | Manual       | 21.1 km/kg |
| 6 | 6     | Nissan Micra Diesel XV          | Jaipur     | 2013 | 86999             | Diesel    | Manual       | 23.08 kmpl |



In [31]:

```
1 feature_matrix = a.iloc[:,0:13]
2 target_vector = a.iloc[:, -1]
```

In [32]:

```
1 print("The feature_matrix has %d rows and %d columns"%(feature_matrix.shape))
```

The feature\_matrix has 7251 rows and 13 columns

In [33]:

```
1 print("The target_vector has %d rows and %d columns"%(np.array(target_vector).resha
```

The target\_vector has 7251 rows and 1 columns

In [45]:

```
1 feature_matrix_standardized = StandardScaler().fit_transform(feature_matrix)
```

```

-----
--
ValueError                                Traceback (most recent call las
t)
~\AppData\Local\Temp\ipykernel_8628\2404365138.py in ?()
----> 1 feature_matrix_standardized = StandardScaler().fit_transform(featur
e_matrix)

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils
\_set_output.py in ?(self, X, *args, **kwargs)
    138     @wraps(f)
    139     def wrapped(self, X, *args, **kwargs):
--> 140         data_to_wrap = f(self, X, *args, **kwargs)
    141         if isinstance(data_to_wrap, tuple):
    142             # only wrap the first output for cross decomposition
    143             return (

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\base.
py in ?(self, X, y, **fit_params)
    874         # non-optimized default implementation; override when a b
etter
    875         # method is possible for a given clustering algorithm
    876         if y is None:
    877             # fit method of arity 1 (unsupervised transformation)
--> 878             return self.fit(X, **fit_params).transform(X)
    879         else:
    880             # fit method of arity 2 (supervised transformation)
    881             return self.fit(X, y, **fit_params).transform(X)

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\prepr
ocessing\_data.py in ?(self, X, y, sample_weight)
    820         Fitted scaler.
    821         """
    822         # Reset internal state before fitting
    823         self._reset()
--> 824         return self.partial_fit(X, y, sample_weight)

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\prepr
ocessing\_data.py in ?(self, X, y, sample_weight)
    857         """
    858         self._validate_params()
    859
    860         first_call = not hasattr(self, "n_samples_seen")
--> 861         X = self._validate_data(
    862             X,
    863             accept_sparse=("csr", "csc"),
    864             dtype=FLOAT_DTYPES,

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\base.
py in ?(self, X, y, reset, validate_separately, **check_params)
    561
    562         if no_val_X and no_val_y:
    563             raise ValueError("Validation should be done on X, y o
r both.")
    564         elif not no_val_X and no_val_y:
--> 565             X = check_array(X, input_name="X", **check_params)
    566             out = X
    567         elif no_val_X and not no_val_y:
    568             y = _check_y(y, **check_params)

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils

```



```

validation.py in ?(array, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, estimator, input_name)
    876         )
    877         array = xp.astype(array, dtype, copy=False)
    878     else:
    879         array = _asarray_with_order(array, order=order, dtype=dtype, xp=xp)
--> 880     except ComplexWarning as complex_warning:
    881         raise ValueError(
    882             "Complex data not supported\n{}\n".format(array)
    883         ) from complex_warning

```

```

~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\_array_api.py in ?(array, dtype, order, copy, xp)
    181     if xp is None:
    182         xp, _ = get_namespace(array)
    183     if xp.__name__ in {"numpy", "numpy.array_api"}:
    184         # Use NumPy API to support order
--> 185     array = numpy.asarray(array, order=order, dtype=dtype)
    186     return xp.asarray(array, copy=copy)
    187 else:
    188     return xp.asarray(array, dtype=dtype, copy=copy)

```

```

~\AppData\Local\Programs\Python\Python310\lib\site-packages\pandas\core\generic.py in ?(self, dtype)
    1996     def __array__(self, dtype: npt.DTypeLike | None = None) -> np.ndarray:
    1997         values = self._values
-> 1998         arr = np.asarray(values, dtype=dtype)
    1999         if (
    2000             astype_is_view(values.dtype, arr.dtype)
    2001             and using_copy_on_write()

```

**ValueError:** could not convert string to float: 'Honda Jazz V'

In [46]:

```
1 algorithm = LogisticRegression(penalty='l2', dual=False, tol=1e-4, C=1.0, fit_intercept=True)
```

In [47]:

```
1 Logistic_Regression_Model = algorithm.fit(feature_matrix_standardized, target_vector)
```

**NameError**

Traceback (most recent call last)

t)

Cell In[47], line 1

```
----> 1 Logistic_Regression_Model = algorithm.fit(feature_matrix_standardized, target_vector)
```

**NameError:** name 'feature\_matrix\_standardized' is not defined

In [48]:

```
1 observation = [[1,0,0.9,0.8000,9.789,9.086,0.956,0.0005,0.675,0.846,0.82541,0.94673
```

In [49]:

```
1 predictions = Logistic_Regression_Model.predict(observation)
2 print('The Model predicted The observation to belong to class %s'%(predictions))
```

```
-----
--
NameError                                Traceback (most recent call las
t)
Cell In[49], line 1
----> 1 predictions = Logistic_Regression_Model.predict(observation)
      2 print('The Model predicted The observation to belong to class %
s'%(predictions))
```

**NameError:** name 'Logistic\_Regression\_Model' is not defined

In [50]:

```
1 print('Algorithm was Trained To predict one of the two classes : %s'%(algorithm.cla
```

```
-----
--
AttributeError                            Traceback (most recent call las
t)
Cell In[50], line 1
----> 1 print('Algorithm was Trained To predict one of the two classes :
%s'%(algorithm.classes_))
```

**AttributeError:** 'LogisticRegression' object has no attribute 'classes\_'

In [51]:

```

1 print("""The Model ssays the probability of the observation we passed Belonging To class['No'] Is %s"""%(algorithm.predict_proba(observation)[0][0]))
2 print()
3 print("""The Model ssays the probability of the observation we passed Belonging To class['Yes'] Is %s"""%(algorithm.predict_proba(observation)[0][1]))

```

```

-----
--
NotFittedError                                Traceback (most recent call last)
Cell In[51], line 1
----> 1 print("""The Model ssays the probability of the observation we passed Belonging To class['No'] Is %s"""%(algorithm.predict_proba(observation)[0][0]))
      2 print()
      3 print("""The Model ssays the probability of the observation we passed Belonging To class['Yes'] Is %s"""%(algorithm.predict_proba(observation)[0][1]))

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear\_model\\_logistic.py:1362, in LogisticRegression.predict\_proba(self, X)

```

1336 def predict_proba(self, X):
1337     """
1338     Probability estimates.
1339     (...)
1360     where classes are ordered as they are in `self.classes_`
1361     """
-> 1362     check_is_fitted(self)
1364     ovr = self.multi_class in ["ovr", "warn"] or (
1365         self.multi_class == "auto"
1366         and (
1367             (...)
1369         )
1370     )
1371     if ovr:

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1390, in check\_is\_fitted(estimator, attributes, msg, all\_or\_any)

```

1385     fitted = [
1386         v for v in vars(estimator) if v.endswith("_") and not v.startswith("_")
1387     ]
1389     if not fitted:
-> 1390         raise NotFittedError(msg % {"name": type(estimator).__name__})

```

**NotFittedError:** This LogisticRegression instance is not fitted yet. Call 'fit' with appropriate arguments before using this estimator.

In [ ]:

|   |  |
|---|--|
| 1 |  |
|---|--|